

AutoML Agent Interface

Select the dataset type

tabular



Choose a dataset

Iris



Select a GROQ LLM Model

gemini-2.0-flash



Run AutoML Agent

Generated Configuration Space Code

```
from ConfigSpace import ConfigurationSpace, Categorical, Float, Integer, Forbidden

def get_configspace():
    cs = ConfigurationSpace()

    # Define hyperparameters for a RandomForestClassifier
    n_estimators = Integer("n_estimators", bounds=(10, 200), default=100)
    max_depth = Integer("max_depth", bounds=(2, 10), default=6)
    min_samples_split = Integer("min_samples_split", bounds=(2, 10), default=2)
    min_samples_leaf = Integer("min_samples_leaf", bounds=(1, 10), default=1)

    cs.add_hyperparameters([n_estimators, max_depth, min_samples_split, min_sample

    # Example of a forbidden clause
    forbidden_clause = ForbiddenAndConjunction(
        ForbiddenEqualsClause(max_depth, 2),
        ForbiddenEqualsClause(min_samples_split, 2)
    )

    cs.add_forbidden_clause(forbidden_clause)

    return cs
```

Generated Scenario Code

```
from smac.scenario import Scenario
from ConfigSpace import ConfigurationSpace

def generate_scenario(cs):
    scenario = Scenario(
        configspace=cs,
        name="gemini-2.0-flashiris20250607_112729",
        output_directory="./automl_results",
        deterministic=False,
        n_workers=4,
        min_budget=1,
        max_budget=10,
        n_trials=10
    )
    return scenario
```

Generated Training Function Code

```
from typing import Any
from ConfigSpace import Configuration
import numpy as np
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
import warnings

def train(cfg: Configuration, dataset: Any, seed: int) -> float:
    """
    Trains a RandomForestClassifier on the given dataset using the provided config

    Args:
        cfg (Configuration): A ConfigSpace Configuration object containing hyperpa
        dataset (Any): A dictionary containing the training data, with 'X' for fea
        seed (int): Random seed for reproducibility.

    Returns:
        float: The average training loss (negative cross-validation score) over 10
    """
    X = dataset['X']
```

```
y = dataset['y']

# Data Preprocessing: StandardScaler
scaler = StandardScaler()

# Model: RandomForestClassifier
model = RandomForestClassifier(
    n_estimators=cfg.get("n_estimators"),
    max_depth=cfg.get("max_depth"),
    min_samples_split=cfg.get("min_samples_split"),
    min_samples_leaf=cfg.get("min_samples_leaf"),
    random_state=seed,
    n_jobs=1 # Explicitly set n_jobs to 1 for consistency
)
pipeline = Pipeline([('scaler', scaler), ('model', model)])

# Cross-validation (10-fold)
with warnings.catch_warnings():
    warnings.filterwarnings("ignore")
    cv_scores = cross_val_score(pipeline, X, y, cv=10, scoring='neg_log_loss',

# Return average loss
loss = -np.mean(cv_scores) # Convert negative log-loss to positive loss

return loss
```

AutoML Agent setup complete!

Loss Value

0.16374038282213835

Prompts Used

▼ {

```
"config" :
```

```
"**TASK**

**Goal:**
Write a Python function named `get_configspace()` that returns a valid
`ConfigurationSpace` object for a classification task.

---

**STRICT OUTPUT RULES**

* Output only the complete `get_configspace()` function and required
imports.
* Do not include any explanations, comments, docstrings, or extra text.
* The code must be syntactically correct, executable, and compatible
with SMAC.

---

**ALLOWED CLASSES**

**Core**

* `ConfigurationSpace`
* `Categorical`
* `Float`
* `Integer`
* `Constant`

**Conditions**

* `EqualsCondition`
* `InCondition`
* `OrConjunction`

**Forbidden Clauses**

* `ForbiddenEqualsClause`
* `ForbiddenAndConjunction` *(must include at least one)*

**Distributions (only if needed)**

* `Beta`
* `Normal`

**Serialization (only if needed)**

* `to_yaml()`
```

```

* `from_yaml()`

---

**CONSTRAINTS**

* Must include at least one `ForbiddenAndConjunction` to block invalid
hyperparameter combinations.

---

**DATASET DESCRIPTION**

* Use the following information to design the configuration space:
  `This is a tabular dataset.
  It has 150 samples and 4 features.
  Feature columns and types:
  - 0: float64
  - 1: float64
  - 2: float64
  - 3: float64

  Feature statistical summary:

  count    150.000000    150.000000    150.000000    150.000000
  mean      5.843333      3.057333      3.758000      1.199333
  std       0.828066      0.435866      1.765298      0.762238
  min       4.300000      2.000000      1.000000      0.100000
  25%       5.100000      2.800000      1.600000      0.300000
  50%       5.800000      3.000000      4.350000      1.300000
  75%       6.400000      3.300000      5.100000      1.800000
  max       7.900000      4.400000      6.900000      2.500000

  Label distribution:
  0      50
  1      50
  2      50
  Name: count, dtype: int64`
  * Hyperparameter choices and model types must be suitable for this
  classification dataset.

---

**SUGGESTED PARAMETERS From OpenML**

* Here are some parameter configurations for this dataset from OpenML (for
inspiration only, not mandatory):

```

```
`[{9551: OpenML Parameter
=====
ID.....: 9551
Flow ID.....: 1068
Flow Name.....: weka.J48(28)_C
Flow URL.....: https://www.openml.org/f/1068
Parameter Name: C
  |__Data Type: option
  |__Default...: 0.25
  |__Value.....: 0.25, 9552: OpenML Parameter
=====
ID.....: 9552
Flow ID.....: 1068
Flow Name.....: weka.J48(28)_M
Flow URL.....: https://www.openml.org/f/1068
Parameter Name: M
  |__Data Type: option
  |__Default...: 2
  |__Value.....: 2}, {9566: OpenML Parameter
=====
ID.....: 9566
Flow ID.....: 1070
Flow Name.....: weka.Ridor(3)_F
Flow URL.....: https://www.openml.org/f/1070
Parameter Name: F
  |__Data Type: option
  |__Default...: 3
  |__Value.....: 3, 9567: OpenML Parameter
=====
ID.....: 9567
Flow ID.....: 1070
Flow Name.....: weka.Ridor(3)_S
Flow URL.....: https://www.openml.org/f/1070
Parameter Name: S
  |__Data Type: option
  |__Default...: 1
  |__Value.....: 1, 9570: OpenML Parameter
=====
ID.....: 9570
Flow ID.....: 1070
Flow Name.....: weka.Ridor(3)_N
Flow URL.....: https://www.openml.org/f/1070
Parameter Name: N
  |__Data Type: option
  |__Default...: 2.0
  |__Value.....: 2.0}]`
```

****IMPORTANT RULE****

* Do ****not**** use any class, function, method, or module outside the ****ALLOWED CLASSES**** list.

****EXAMPLES****

* See provided examples for valid usage of hyperparameters, conditions, forbidden clauses, and priors.

[EXAMPLES]

Example 1: Basic ConfigurationSpace

```
```python
from ConfigSpace import ConfigurationSpace
```

```
cs = ConfigurationSpace(
 space={
 "C": (-1.0, 1.0),
 "max_iter": (10, 100),
 },
)
```
```

Example 2: Adding Hyperparameters

```
```python
from ConfigSpace import ConfigurationSpace, Categorical, Float, Integer
```

```
kernel_type = Categorical('kernel_type', ['linear', 'poly', 'rbf', 'sigmoid'])
degree = Integer('degree', bounds=(2, 4), default=2)
coef0 = Float('coef0', bounds=(0, 1), default=0.0)
gamma = Float('gamma', bounds=(1e-5, 1e2), default=1, log=True)
```

```
cs = ConfigurationSpace()
cs.add([kernel_type, degree, coef0, gamma])
```
```

Example 3: Adding Conditions

```
```python
from ConfigSpace import EqualsCondition, InCondition, OrConjunction
```

```
cond_1 = EqualsCondition(degree, kernel_type, 'poly')
cond_2 = OrConjunction(
 EqualsCondition(coef0, kernel_type, 'poly'),
 EqualsCondition(coef0, kernel_type, 'sigmoid')
```



```

 EqualsCondition(eccs, kernel_type, ['sigmoid',
)
cond_3 = InCondition(gamma, kernel_type, ['rbf', 'poly', 'sigmoid'])
'''

Example 4: Adding Forbidden Clauses
'''python
from ConfigSpace import ForbiddenEqualsClause, ForbiddenAndConjunction

penalty_and_loss = ForbiddenAndConjunction(
 ForbiddenEqualsClause(penalty, "l1"),
 ForbiddenEqualsClause(loss, "hinge")
)
constant_penalty_and_loss = ForbiddenAndConjunction(
 ForbiddenEqualsClause(dual, "False"),
 ForbiddenEqualsClause(penalty, "l2"),
 ForbiddenEqualsClause(loss, "hinge")
)
penalty_and_dual = ForbiddenAndConjunction(
 ForbiddenEqualsClause(dual, "False"),
 ForbiddenEqualsClause(penalty, "l1")
)
'''

Example 5: Serialization
'''python
from pathlib import Path
from ConfigSpace import ConfigurationSpace

path = Path("configspace.yaml")
cs = ConfigurationSpace(
 space={
 "C": (-1.0, 1.0),
 "max_iter": (10, 100),
 },
)
cs.to_yaml(path)
loaded_cs = ConfigurationSpace.from_yaml(path)
'''

Example 6: Priors
'''python
import numpy as np
from ConfigSpace import ConfigurationSpace, Float, Categorical, Beta, Normal

cs = ConfigurationSpace(
 space={
 "lr": Float(
 'lr',
 bounds=(1e-5, 1e-1),
 default=1e-2

```

```
 default=1e-5,
 log=True,
 distribution=Normal(1e-3, 1e-1)
),
 "dropout": Float(
 'dropout',
 bounds=(0, 0.99),
 default=0.25,
 distribution=Beta(alpha=2, beta=4)
),
 "activation": Categorical(
 'activation',
 items=['tanh', 'relu'],
 weights=[0.2, 0.8]
),
}
)"
```

"scenario" :

```
"""
```

```
Objective:
```

Generate a **Python function** named `generate_scenario(cs)` that returns a valid `Scenario` object configured for SMAC (v2.0+), strictly following the rules below.

```

```

```
Output Format Rules (Strict):
```

- \* Output **only** the function `generate_scenario(cs)` and the **necessary** import statements.
- \* Use **Python 3.10 syntax** but **do not** include type annotations for the function or parameters.
- \* The code must be **fully executable** with the latest **SMAC v2.0+** version.
- \* Output **only valid Python code** – **no comments**, **no explanations**, **no extra text**, and **no example usage**.
- \* The function must be **self-contained**.

```

```

```
Functional Requirements:
```

- \* The input `cs` is a `ConfigurationSpace` object.
- \* Return a `Scenario` configured with the following:
  - \* `name`: `"gemini-2.0-flashiris20250607_112729"`
  - \* `output_directory`: `"./automl_results"`
  - \* `deterministic`: `False` (enable variability)
  - \* `n_workers`: greater than 1 (to enable parallel optimization)
  - \* `min_budget` and `max_budget`: set appropriately for multi-fidelity tuning (e.g., training epochs)
  - \* `n_trials`: 10

```

```

```
Other Parameters to Consider:
```

```
configspace : ConfigurationSpace
```

The configuration space from which to sample the configurations.

```
name : str | None, defaults to None
```

The name of the run. If no name is passed, SMAC generates a hash from the meta data.

Specify this argument to identify your run easily.

```
output_directory : Path, defaults to Path("smac3_output")
```

The directory in which to save the output. The files are saved in `./output_directory/name/seed`.

```
deterministic : bool, defaults to False
```

If deterministic is set to true, only one seed is passed to the target

function.

Otherwise, multiple seeds (if `n_seeds` of the intensifier is greater than 1) are passed

to the target function to ensure generalization.

`objectives` : str | list[str] | None, defaults to "cost"

The objective(s) to optimize. This argument is required for multi-objective optimization.

`crash_cost` : float | list[float], defaults to `np.inf`

Defines the cost for a failed trial. In case of multi-objective, each objective can be associated with

a different cost.

`termination_cost_threshold` : float | list[float], defaults to `np.inf`

Defines a cost threshold when the optimization should stop. In case of multi-objective, each objective *must* be

associated with a cost. The optimization stops when all objectives crossed the threshold.

`walltime_limit` : float, defaults to `np.inf`

The maximum time in seconds that SMAC is allowed to run.

`cputime_limit` : float, defaults to `np.inf`

The maximum CPU time in seconds that SMAC is allowed to run.

`trial_walltime_limit` : float | None, defaults to None

The maximum time in seconds that a trial is allowed to run. If not specified,

no constraints are enforced. Otherwise, the process will be spawned by `pynisher`.

`trial_memory_limit` : int | None, defaults to None

The maximum memory in MB that a trial is allowed to use. If not specified,

no constraints are enforced. Otherwise, the process will be spawned by `pynisher`.

`n_trials` : int, defaults to 100

The maximum number of trials (combination of configuration, seed, budget, and instance, depending on the task) to run.

`use_default_config`: bool, defaults to False.

If True, the configspace's default configuration is evaluated in the initial design.

For historic benchmark reasons, this is False by default.

Notice, that this will result in `n_configs + 1` for the initial design. Respecting `n_trials`,

this will result in one fewer evaluated configuration in the optimization.

`instances` : list[str] | None, defaults to None

Names of the instances to use. If None, no instances are used.

Instances could be dataset names, seeds, subsets, etc.

`instance_features` : dict[str, list[float]] | None, defaults to None

Instances can be associated with features. For example, meta data of

```

the dataset (mean, var, ...) can be
 incorporated which are then further used to expand the training data of
the surrogate model.
 min_budget : float | int | None, defaults to None
 The minimum budget (epochs, subset size, number of instances, ...) that
is used for the optimization.
 Use this argument if you use multi-fidelity or instance optimization.
 max_budget : float | int | None, defaults to None
 The maximum budget (epochs, subset size, number of instances, ...) that
is used for the optimization.
 Use this argument if you use multi-fidelity or instance optimization.
 seed : int, defaults to 0
 The seed is used to make results reproducible. If seed is -1, SMAC will
generate a random seed.
 n_workers : int, defaults to 1
 The number of workers to use for parallelization. If `n_workers` is
greater than 1, SMAC will use
 Dask to parallelize the optimization.

Reminder: The output must be limited to:

* Valid `import` statements
* A single `generate_scenario(cs)` function that returns a properly configured
`Scenario` object

Based on the following SMAC documentation, analyze the dataset
characteristics and choose appropriate:
 1. Facade type (e.g., MultiFidelityFacade for multi-fidelity
optimization)
 2. Budget settings (min_budget and max_budget)
 3. Number of workers (n_workers)
 4. Other relevant scenario parameters

SMAC Documentation:
Getting Started

#
SMAC needs four core components (configuration space, target function, scenario
and a facade) to run an
optimization process, all of which are explained on this page.
They interact in the following way:
Interaction of SMAC's components
Configuration Space

#
The configuration space defines the search space of the hyperparameters and,
therefore, the tunable parameters' legal

```

Therefore, the search parameters range

ranges and default values.

from

ConfigSpace

import

ConfigSpace

cs

=

ConfigurationSpace

({

"myfloat"

:

(

0.1

,

1.5

),

# Uniform Float

"myint"

:

(

2

,

10

),

# Uniform Integer

"species"

:

[

"mouse"

,

"cat"

,

"dog"

],

# Categorical

})

Please see the documentation of

ConfigurationSpace

for more details.

Target Function

#

The target function takes a configuration from the configuration space and returns a performance value.

For example, you could use a Neural Network to predict on your data and get some validation performance.

If, for instance, you would tune the learning rate of the Network's optimizer, every learning rate will

every learning rate will

change the final validation performance of the network. This is the target function.

SMAC tries to find the best performing learning rate by trying different values and evaluating the target function - in an efficient way.

```
def
train
(
self
,
config
:
Configuration
,
seed
:
int
)
->
float
:
model
=
MultiLayerPerceptron
(
learning_rate
=
config
[
"learning_rate"
])
model
.
fit
(
...
)
accuracy
=
model
.
validate
(
...
)
return
1
```



```
└
-
accuracy
SMAC always minimizes (the smaller the better)
Note
In general, the arguments of the target function depend on the intensifier.
However,
in all cases, the first argument must be the configuration (arbitrary argument
name is possible here) and a seed.
If you specified instances in the scenario, SMAC requires
instance
as argument additionally. If you use
SuccessiveHalving
or
Hyperband
as intensifier but you did not specify instances, SMAC passes
budget
as
argument to the target function. But don't worry: SMAC will tell you if
something is missing or if something is not
used.
Warning
SMAC
always
minimizes the value returned from the target function.
Warning
SMAC passes either
instance
or
budget
to the target function but never both.
Scenario
#
The
Scenario
is used to provide environment variables. For example,
if you want to limit the optimization process by a time limit or want to
specify where to save the results.
from
smac
import
Scenario
scenario
=
Scenario
(
configspace
```

```
=
CS
,
name
=
"experiment_name"
,
output_directory
=
Path
(
"your_output_directory"
)
walltime_limit
=
120
,
Limit to two minutes
n_trials
=
500
,
Evaluated max 500 trials
n_workers
=
8
,
Use eight workers
...
)
Note
If no
name
is given, a hash of the experiment is used. Running the same experiment again
at a later time will result in exactly the same hash. This is important,
because the optimization will warmstart on the preexisting evaluations, if not
otherwise specified in the
Facade
.
Facade
#
Warn
By default Facades will try to warmstart on preexisting logs. This behavior can
be specified using the
overwrite
parameter.
A
```

facade

is the entry point to SMAC, which constructs a default optimization pipeline for you. SMAC offers various facades, which satisfy many common use cases and are crucial to achieving peak performance. The idea behind the facades is to provide a simple interface to all of SMAC's components, which is easy to use and understand and without the need of deep diving into the material. However, experts are invited to change the components to their specific hyperparameter optimization needs. The following table (horizontally scrollable) shows you what is supported and reveals the default components

:

Black-Box

Hyperparameter Optimization

Multi-Fidelity

Algorithm Configuration

Random

Hyperband

#Parameters

low

low/medium/high

low/medium/high

low/medium/high

low/medium/high

low/medium/high

Supports Instances

✗

✓

✓

✓

✓

✓

Supports Multi-Fidelity

✗

✗

✓

✓

✗

✓

Initial Design

Sobol

Sobol

Random

Default

Default

Default  
Surrogate Model  
Gaussian Process  
Random Forest  
Random Forest  
Random Forest  
Not used  
Not used  
Acquisition Function  
Expected Improvement  
Log Expected Improvement  
Log Expected Improvement  
Expected Improvement  
Not used  
Not used  
Acquisition Maximizer  
Local and Sorted Random Search  
Local and Sorted Random Search  
Local and Sorted Random Search  
Local and Sorted Random Search  
Not Used  
Not Used  
Intensifier  
Default  
Default  
Hyperband  
Default  
Default  
Hyperband  
Runhistory Encoder  
Default  
Log  
Log  
Default  
Default  
Default  
Random Design Probability  
8.5%  
20%  
20%  
50%  
Not used  
Not used  
Info  
The multi-fidelity facade is the closest implementation to  
BOHB  
.

## Note

We want to emphasize that SMAC is a highly modular optimization framework. The facade accepts many arguments to specify components of the pipeline. Please also note, that in contrast to previous versions, instantiated objects are passed instead of kwargs

.

The facades can be imported directly from the

smac

module.

from

smac

import

BlackBoxFacade

as

BBFacade

from

smac

import

HyperparameterOptimizationFacade

as

HPOFacade

from

smac

import

MultiFidelityFacade

as

MFFacade

from

smac

import

AlgorithmConfigurationFacade

as

ACFacade

from

smac

import

RandomFacade

as

RFacade

from

smac

import

HyperbandFacade

as

HBFacade

smac

```
=
HP0Facade
(
scenario
=
scenario
,
target_function
=
train
)
smac
=
MFFacade
(
scenario
=
scenario
,
target_function
=
train
)
smac
=
ACFacade
(
scenario
=
scenario
,
target_function
=
train
)
smac
=
RFacade
(
scenario
=
scenario
,
target_function
=
train
)
)
```

```
smac
=
HBFacade
(
scenario
=
scenario
,
target_function
=
train
)
```

## Multi-Fidelity Optimization

```
#
Multi-fidelity refers to running an algorithm on multiple budgets (such as
number of epochs or
subsets of data) and thereby evaluating the performance prematurely. You can
run a multi-fidelity optimization
when using
Successive Halving
or
Hyperband
.
Hyperband
is the default intensifier in the
multi-fidelity facade
and requires the arguments
min_budget
and
max_budget
in the scenario if no instances are used.
In general, multi-fidelity works for both real-valued and instance budgets. In
the real-valued case,
the budget is directly passed to the target function. In the instance case, the
budget is not passed to the
target function but
min_budget
and
max_budget
are used internally to determine the number of instances of
each stage. That's also the reason why
min_budget
and
max_budget
are
not required
```

when using instances:

The

`max_budget`

is simply the max number of instances, whereas the

`min_budget`

is simply 1.

Warning

`smac.main.config_selector.ConfigSelector`

contains the

`min_trials`

parameter. This parameter determines

how many samples are required to train the surrogate model. If budgets are

involved, the highest budgets

are checked first. For example, if `min_trials` is three, but we find only two trials in the runhistory for

the highest budget, we will use trials of a lower budget instead.

Please have a look into our

[multi-fidelity examples](#)

to see how to use

multi-fidelity optimization in real-world applications.

Components

#

In addition to the basic components mentioned in

[Getting Started](#)

, all other components are

explained in the following paragraphs to give a better picture of SMAC. These components are all used to guide

the optimization process and simple changes can influence the results drastically.

Before diving into the components, we shortly want to explain the main Bayesian optimization loop in SMAC.

The

`SMBO`

receives all instantiated components from the facade and the logic happens here.

In general, a while loop is used to ask for the next trial, submit it to the runner, and wait for the runner to

finish the evaluation. Since the runner and the

`SMBO`

object are decoupled, the while loop continues and asks for even

more trials (e.g., in case of multi-threading), which also can be submitted to the runner. If all workers are

occupied, SMAC will wait until a new worker is available again. Moreover, limitations like wallclock time and remaining

trials are checked in every iteration.

Surrogate Model



Surrogate Model

#

The surrogate model is used to approximate the objective function of configurations. In previous versions, the model was referred to as the Empirical Performance Model (EPM). Mostly, Bayesian optimization is used/associated with Gaussian processes. However, SMAC also incorporates random forests as surrogate models, which makes it possible to optimize for higher dimensional and complex spaces.

The data used to train the surrogate model is collected by the runhistory encoder (receives data from the runhistory and transforms it). If budgets are involved, the highest budget which satisfies min\_trials

(defaults to 1) in

smac.main.config\_selector

is

used. If no budgets are used, all observations are used.

If you are using instances, it is recommended to use instance features. The model is trained on each instance

associated with its features. Imagine you have two hyperparameters, two instances and no instance features, the model would be trained on:

HP 1

HP 2

Objective Value

0.1

0.8

0.5

0.1

0.8

0.75

505

7

2.4

505

7

1.3

You can see that the same inputs lead to different objective values because of two instances. If you associate

each instance with a feature, you would end-up with the following data points:

HP 1

HP 2

Instance Feature

Objective Value

0.1

0.8

0

0

0.5

0.1

0.8

1

0.75

505

7

0

2.4

505

7

1

1.3

The steps to receiving data are as follows:

The intensifier requests new configurations via

`next(self.config_generator)`

.

The config selector collects the data via the runhistory encoder which iterates over the runhistory trials.

The runhistory encoder only collects trials which are in `considered_states`

and timeout trials. Also, only the

highest budget is considered if budgets are used. In this step, multi-objective values are scalarized using the `normalize_costs`

function (uses

`objective_bounds`

from the runhistory) and the multi-objective algorithm.

For example, when ParEGO is used, the scalarization would be different in each training.

The selected trial objectives are transformed (e.g., log-transformed, depending on the selected

encoder).

The hyperparameters might still have inactive values. The model takes care of that after the collected data

are passed to the model.

Acquisition Function

#

Acquisition functions are mathematical techniques that guide how the parameter space should be explored during Bayesian optimization. They use the predicted mean and predicted variance generated by the surrogate model.

The acquisition function is used by the acquisition maximizer (see next section). Otherwise, SMAC provides

a bunch of different acquisition functions (Lower Confidence Bound, Expected Improvement, Probability Improvement,

inompson, integrated acquisition functions and prior acquisition functions). We refer to literature

for more information about acquisition functions.

#### Note

The acquisition function calculates the acquisition value for each configuration. However, the configurations are provided by the acquisition maximizer. Therefore, the acquisition maximizer is responsible for receiving the next configurations.

#### Acquisition Maximize

#

The acquisition maximizer is a wrapper for the acquisition function. It returns the next configurations. SMAC

supports local search, (sorted) random search, local and (sorted) random search, and differential evolution.

While local search checks neighbours of the best configurations, random search makes sure to explore the configuration space. When using sorted random search, random configurations are sorted by the value of the acquisition function.

#### Warning

Pay attention to the number of challengers: If you experience RAM issues or long computational times in the acquisition function, you might lower the number of challengers.

The acquisition maximizer also incorporates the

#### Random Design

. Please see the

#### ChallengerList

for more information.

#### Initial Design

#

The surrogate model needs data to be trained. Therefore, the initial design is used to generate the initial data points.

We provide random, latin hypercube, sobol, factorial and default initial designs. The default initial design uses

the default configuration from the configuration space and with the factorial initial design, we generate corner

points of the configuration space. The sobol sequences are an example of quasi-random low-discrepancy sequences and

the latin hypercube design is a statistical method for generating a near-random sample of parameter values from a multidimensional distribution.

The initial design configurations are yielded by the config selector first.

Moreover, the config selector keeps

track of which configurations already have been returned to make sure a configuration is not returned twice.

#### Random Design

#

The random design is used in the acquisition maximizer to tell whether the next configuration should be random or sampled from the acquisition function. For example, if we use a random design with a probability of 50%, we have a 50% chance to sample a random configuration and a 50% chance to sample a configuration from the acquisition function (although the acquisition function includes exploration and exploitation trade-off already). This design makes sure that the optimization process is not stuck in a local optimum and we are guaranteed to find the best configuration over time. In addition to simple probability random design, we also provide annealing and modulus random design.

Intensifier

#

The intensifier compares different configurations based on evaluated :term: trial<Trial>

so far. It decides

which configuration should be intensified

or, in other words, if a configuration is worth to spend more time on (e.g., evaluate another seed pair, evaluate on another instance, or evaluate on a higher budget).

Warning

Always pay attention to

max\_config\_calls

or

n\_seeds

: If this argument is set high, the intensifier might spend a lot of time on a single configuration.

Depending on the components and arguments, the intensifier tells you which seeds, budgets, and/or instances

are used throughout the optimization process. You can use the methods uses\_seeds

,

uses\_budgets

, and

uses\_instances

(directly callable via the facade) to (sanity-)check whether the intensifier uses these arguments.

Another important fact is that the intensifier keeps track of the current incumbent (a.k.a. the best configuration

found so far). In case of multi-objective, multiple incumbents could be found.

All intensifiers support multi-objective, multi-fidelity, and multi-threading:

Multi-Objective: Keeping track of multiple incumbents at once.

Multi-Fidelity: Incorporating instances or budgets.

Multi-Threading: Intensifier are implemented as generators so that calling next

on the intensifier can be

repeated as often as needed. Intensifier are not required to receive results as the results are directly taken from

the runhistory.

Note

All intensifiers are working on the runhistory and recognize previous logged trials (e.g., if the user already evaluated something beforehand). Previous configurations (in the best case, also complete trials) are added to the queue/tracker again so that they are integrated into the intensification process.

That means continuing a run as well as incorporating user inputs are natively supported.

Configuration Selector

#

The configuration selector uses the initial design, surrogate model, acquisition maximizer/function, runhistory, runhistory encoder, and random design to select the next configuration. The configuration selector is directly used by the intensifier and is called everytime a new configuration is requested.

The idea behind the configuration selector is straight forward:

Yield the initial design configurations.

Train the surrogate model with the data from the runhistory encoder.

Get the next

retrain\_after

configurations from the acquisition function/maximizer and yield them.

After all

retrain\_after

configurations were yield, go back to step 2.

Note

The configuration selector is a generator and yields configurations. Therefore, the current state of the

selector is saved and when the intensifier calls

next

, the selector continues there where it stopped.

Note

Everytime the surrogate model is trained, the multi-objective algorithm is updated via

update\_on\_iteration\_start

.

Multi-Objective Algorithm

#

The multi-objective algorithm is used to scalarize multi-objective values. The

multi-objective algorithm gets normalized objective values passed and returns a single value. The resulting value (called by the runhistory encoder) is then used to train the surrogate model.

#### Warning

Depending on the multi-objective algorithm, the values for the runhistory encoder might differ each time the surrogate model is trained. Let's take ParEGO for example: Everytime a new configuration is sampled (see ConfigSelector), the objective weights are updated. Therefore, the scalarized values are different and the acquisition maximizer might return completely different configurations.

#### RunHistory

#

The runhistory holds all (un-)evaluated trials of the optimization run. You can use the runhistory to get (running) configs, (running) trials, trials of a specific config, and more. The runhistory encoder iterates over the runhistory to receive data for the surrogate model. The following code shows how to iterate over the runhistory:

```
smac
=
HPOFacade
(
...
)
Iterate over all trials
for
trial_info
,
trial_value
in
smac
.
runhistory
.
items
():
Trial info
config
=
trial_info
.
config
instance
=
trial_info
```

```
.
instance
budget
=
trial_info
.
budget
seed
=
trial_info
.
seed
Trial value
cost
=
trial_value
.
cost
time
=
trial_value
.
time
status
=
trial_value
.
status
starttime
=
trial_value
.
starttime
endtime
=
trial_value
.
endtime
additional_info
=
trial_value
.
additional_info
Iterate over all configs
for
config
in
```

```
smac
.
runhistory
.
get_configs
():
Get the cost of all trials of this config
average_cost
=
smac
.
runhistory
.
average_cost
(
config
)
Warning
The intensifier uses a callback to update the incumbent everytime a new trial
is added to the runhistory.
RunHistory Encoder
#
The runhistory encoder is used to encode the runhistory data into a format that
can be used by the surrogate model.
Only trials with the status
considered_states
and timeout trials are considered. Multi-objective values are
scalarized using the
normalize_costs
function (uses
objective_bounds
from the runhistory). Afterwards, the
normalized value is processed by the multi-objective algorithm.
Callback
#
Callbacks provide the ability to easily execute code before, inside, and after
the Bayesian optimization loop.
To add a callback, you have to inherit from
smac.Callback
and overwrite the methods (if needed).
Afterwards, you can pass the callbacks to any facade.
from
smac
import
MultiFidelityFacade
,
Callback
```



```
class
CustomCallback
(
Callback
):
def
on_start
(
self
,
smbo
:
SMBO
)
->
None
:
pass
def
on_end
(
self
,
smbo
:
SMBO
)
->
None
:
pass
def
on_iteration_start
(
self
,
smbo
:
SMBO
)
->
None
:
pass
def
on_iteration_end
(
```

```

\
self
,
smbo
:
SMBO
,
info
:
RunInfo
,
value
:
RunValue
)
->
bool
|
None
:
We just do a simple printing here
print
(
info
,
value
)
smac
=
MultiFidelityFacade
(
...
callbacks
=
[
CustomCallback
()]
)
smac
.
optimize
()
```

Please analyze the dataset and documentation to determine:

1. Should multi-fidelity optimization be used? (Consider dataset size and training time)
2. What budget range is appropriate? (Consider training epochs or data subsets)
3. How many workers should be used? (Consider available resources)

3. How many workers should be used? (consider available resources)
4. Are there any special considerations for this dataset type?

Then generate a scenario configuration that best suits this dataset.

"

```
"train_function" :
```

**\*\*Generate production-grade Python code for a machine learning training function with the following STRICT requirements:\*\***

---

**### \*\*Function signature\*\* must be:**

```
```python
from ConfigSpace import Configuration
def train(cfg: Configuration, dataset: Any, seed: int) -> float:
    ...
```

**Function Behavior Requirements:**

* The function **must accept** a ``dataset`` dictionary with:

- * ``dataset['X']``: feature matrix or input tensor
- * ``dataset['y']``: label vector or label tensor

* Assume ``cfg`` is a sampled configuration object:

* Access primitive values using ``cfg.get('key')`` (only ``int``, ``float``, ``str``, etc.).

* **Do not access or manipulate non-primitive hyperparameter objects**.

* The function must return the **average training loss** over 10 epochs.

* You must carefully read and follow the dataset description provided, which includes:

- * Data format and dimensions
- * Required preprocessing steps
- * Special handling requirements
- * Framework-specific considerations

```
```python
return loss # float
...
```

\* Lower ``loss`` means a better model.

---

**### \*\*Frameworks\*\***

You may choose **PyTorch**, **TensorFlow**, or **scikit-learn**, depending on

the dataset and supporting code provided.

---

### \*\*Model Requirements\*\*

- \* Infer input and output dimensions dynamically from the dataset
- \* Follow the data format requirements specified in the dataset description
- \* Handle any necessary data transformations as described in the dataset description

---

### \*\*Supporting Code Provided:\*\*

- \* ConfigSpace definition: ``from ConfigSpace import ConfigurationSpace, Categorical, Float, Integer, ForbiddenAndConjunction, ForbiddenEqualsClause`

```
def get_configspace():
 cs = ConfigurationSpace()

 # Define hyperparameters for a RandomForestClassifier
 n_estimators = Integer("n_estimators", bounds=(10, 200), default=100)
 max_depth = Integer("max_depth", bounds=(2, 10), default=6)
 min_samples_split = Integer("min_samples_split", bounds=(2, 10), default=2)
 min_samples_leaf = Integer("min_samples_leaf", bounds=(1, 10), default=1)

 cs.add_hyperparameters([n_estimators, max_depth, min_samples_split,
min_samples_leaf])

 # Example of a forbidden clause
 forbidden_clause = ForbiddenAndConjunction(
 ForbiddenEqualsClause(max_depth, 2),
 ForbiddenEqualsClause(min_samples_split, 2)
)

 cs.add_forbidden_clause(forbidden_clause)

 return cs
`

* SMAC scenario: `from smac.scenario import Scenario
from ConfigSpace import ConfigurationSpace

def generate_scenario(cs):
 scenario = Scenario(
 configspace=cs,
```

```

 name="gemini-2.0-flashiris20250607_112729",
 output_directory="./automl_results",
 deterministic=False,
 n_workers=4,
 min_budget=1,
 max_budget=10,
 n_trials=10
)
 return scenario

```

\* Dataset description: `This is a tabular dataset.  
It has 150 samples and 4 features.

Feature columns and types:

- 0: float64
- 1: float64
- 2: float64
- 3: float64

Feature statistical summary:

	0	1	2	3
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333
std	0.828066	0.435866	1.765298	0.762238
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

Label distribution:

```

0 50
1 50
2 50

```

Name: count, dtype: int64`

---

### \*\*Additional Instructions\*\*

- \* The code must not hardcode dataset dimensions
- \* The function must be runnable and not assume unavailable classes or modules
- \* You must only output the `def train(...)` function and nothing else
- \* Always check dataset description for format hints and requirements before processing

Data preprocessing involves scaling the features since they have different ranges. Consider StandardScaler or MinMaxScaler.

Feature engineering could involve creating polynomial features or feature interactions, but given the small number of features, this might lead to overfitting.

A common challenge is that the dataset is small, which can lead to overfitting. Regularization techniques are important.

The dataset is balanced, so accuracy is an appropriate metric. However, it's still useful to look at precision, recall, and F1-score, especially if you want to understand class-specific performance.

Consider using cross-validation to evaluate the model performance robustly due to the limited size of the dataset.

When interpreting the model, be mindful of the curse of dimensionality if feature engineering is performed aggressively.

The Iris dataset is linearly separable; therefore, simpler models can achieve high accuracy.

Apply dimensionality reduction techniques such as PCA or LDA to identify the most important features and potentially improve model performance and interpretability.

Investigate non-linear relationships between features and the target variable using techniques like kernel methods or decision tree-based models.

Be cautious about outliers, as they can significantly impact model performance. Consider using robust scaling techniques or outlier removal methods if necessary.

When building classification models, evaluate and compare the performance of different algorithms, such as logistic regression, support vector machines, and decision trees.

Visualize the data using scatter plots or pair plots to gain insights into the relationships between features and the target variable.

Address multicollinearity among features by using techniques like variance inflation factor (VIF) or principal component analysis (PCA)."

}



